

Probing Structural Bias in Stable Diffusion 3: A Stereotype Content Model Analysis

Shamima Hossain

Brac University

Dhaka, Bangladesh

shamima.hossain@g.bracu.ac.bd

Abstract—Modern text-to-image diffusion models act as mirrors of human social cognition, yet most bias audits focus on *frequency* (e.g. gender ratios in generated occupations) rather than on *structure* the shape of the stereotype manifold in latent space. We probe Stable Diffusion 3 Medium (SD3) against the Stereotype Content Model (SCM), social psychology’s two-dimensional theory of group perception along *warmth* and *competence*, using 6,600 generated portraits of 66 social groups drawn from [1] and [2]. We build a probabilistic graphical model in which each group’s image embedding distribution $z_g \sim \mathcal{N}(\mu_g, \Sigma_g)$ is projected onto lexicon-defined warmth/competence axes, and we fit a Bayesian errors-in-variables regression between the model’s coordinates and validated human ratings. SD3 encodes SCM structure to a statistically reliable degree ($\beta_W=0.48$, 94% HDI [0.26, 0.69]; $\beta_C=0.32$, 94% HDI [0.10, 0.55]), but unevenly: per-quadrant Kullback–Leibler divergence shows the *contempt* quadrant is reproduced faithfully (KL=0.68) while the *envy* quadrant the home of high-competence, low-warmth stereotypes such as CEO, lawyer, and financial advisor collapses toward the origin (KL=3.51). The finding reframes diffusion-model bias as a question of stereotype geometry rather than stereotype presence. We release the SCM-aligned portrait corpus and probing code for reuse.

Index Terms—text-to-image models, Stable Diffusion, stereotype content model, probabilistic graphical models, bias auditing, Bayesian regression, representational fairness

I. INTRODUCTION

Text-to-image (T2I) diffusion models are now widely deployed in advertising, education, journalism, and consumer apps; their outputs increasingly shape what whole categories of people are imagined to “look like.” A growing audit literature has shown that these systems reproduce demographic imbalances they generate fewer women when prompted for “CEO,” more Black faces when prompted for “criminal,” and so on [10], [11]. This evidence is critical but operates almost entirely at the level of *marginal frequencies*: how often each attribute appears for each prompt.

A complementary and largely unexamined question is *structural*: when the model places a thousand social groups in its internal representation, what is the shape of the resulting manifold? Is it organized along the same dimensions humans use to judge groups, and if so, where does that organization break down? This is the question we ask.

Social psychology has a strong empirical answer for humans. The *Stereotype Content Model* (SCM) of Fiske, Cuddy, Glick, and Xu [1] demonstrates that group perception across cultures organizes into a two-dimensional topology of **warmth**

(perceived intent: friendly vs. hostile) and **competence** (perceived ability: capable vs. incapable), yielding four quadrants of mixed stereotypes: *Admiration* (warm + competent: e.g. doctors, nurses), *Contempt* (cold + incompetent: e.g. the unhoused, welfare recipients), *Paternalism* (warm + incompetent: e.g. elderly people, housewives), and *Envy* (cold + competent: e.g. CEOs, lawyers, financial advisors). This 2×2 topology has been replicated in over 40 countries and recently recovered in large language models such as GPT-4o [8]. Whether it survives in the *visual* outputs of a diffusion model is open.

Contributions. This paper contributes:

- 1) A probabilistic graphical model framing of T2I representational bias, in which each social group is a latent 2D point in (warmth, competence) space and observed pixels are noisy projections through CLIP/SD encoders.
- 2) A dual-probe methodology on *text* encoders (CLIP-L, OpenCLIP bigG, T5-XXL) and on *image* outputs (CLIP, DINOv2) that localizes where SCM structure lives in the SD3 stack.
- 3) Empirical results showing that SD3 internalizes SCM *partially* and *asymmetrically*: low-status quadrants (contempt) align well with human ratings while the envy quadrant systematically collapses toward the origin.
- 4) A released, reproducible portrait corpus of 6,600 SD3-generated images annotated with ground-truth SCM coordinates, along with all probing code.

The remainder of the paper presents background (§II), methodology and the graphical model (§III), implementation (§IV), results and analysis (§V), usefulness (§VI), and limitations and future work (§VII).

II. BACKGROUND AND RELATED WORK

A. The Stereotype Content Model (SCM)

Fiske et al. [1] introduced the SCM as a unifying account of how perceivers organize their judgments of social groups. Drawing on intergroup-emotion research and the classic agency–communion duality in person perception, they showed across multiple U.S. samples that ratings of 23 societal groups consistently collapse onto two orthogonal factors warmth (*W*) and competence (*C*) and that the resulting 2×2 space cleanly partitions groups into four *mixed-stereotype* quadrants associated with distinct emotional responses (admiration, envy,

pity, and contempt) and distinct behavioral tendencies toward each group [1].

The model has since been replicated in over forty countries [5] and extended along multiple lines. Imhoff and Woelki [3] demonstrated that warmth and competence are read directly from human faces. Nicolas, Bai, and Fiske [4] constructed a validated open-ended SCM dictionary of high-loading adjectives, providing the linguistic anchors used in this paper to define warmth and competence axes in embedding space. He, Caliskan, Houwer, and Spears [2] extended SCM ratings to 60 contemporary U.S. occupations with z -scored warmth and competence means, supplying the occupation half of our ground-truth set.

B. SCM in Modern AI Systems

Caliskan, Bryson, and Narayanan [9] first showed that word-embedding geometry recovers documented human social biases including occupational and racial associations via the WEAT test. The line of work extending these embedding probes from GloVe to contextualized models led naturally to direct SCM probes of language models. Fraser, Kiritchenko, and Nejadgholi [6] mapped LLM completions onto SCM coordinates and found systematic quadrant alignment with human stereotype data; Nicolas and Caliskan [7] formalized *construct differentiability* as a desideratum for stereotype-content auditing. Most recently, a study in *The Gerontologist* [8] showed that GPT-4o’s age-related outputs cluster precisely in the Paternalism quadrant, recovering the SCM topology from a model never trained on SCM data.

For text-to-image models, audits have so far concentrated on *marginal* or *frequency* bias. Bianchi et al. [10] documented that Stable Diffusion amplifies gender and racial stereotypes in occupations beyond U.S. Bureau of Labor Statistics base rates. Ghosh and Caliskan [11] measured how T2I outputs encode nationality stereotypes via face embeddings. Hofmann et al. [12], while focused on language, demonstrated that “aligned” models can retain covertly racist representations that pass surface audits evidence that bias may live in latent geometry rather than in observable token frequencies. Our work asks the analogous question for image diffusion: does the latent geometry exhibit the 2D SCM topology, and if so, with what fidelity?

C. Probabilistic Latent-Variable Framing of Diffusion Models

Diffusion models are themselves deep hierarchical latent-variable models [13], [14]: a prompt y conditions a denoising trajectory $z_T \rightarrow z_{T-1} \rightarrow \dots \rightarrow z_0$ over a sequence of latents. Recent work on probabilistic circuits and continuous latent-variable hierarchies [15] has formalized how structured factorizations sit inside such pipelines. Our contribution is to view the social-perception properties of the *conditional* distribution $P_\theta(z_0 | y)$ through the SCM lens: if the model has internalized human stereotype geometry, then groups sharing a quadrant should occupy spatially coherent regions of latent space.

D. Halo Effects and Visual Stereotyping

A separate but converging line of work shows that visual representation itself carries social-evaluative content. Gulati et al. [16] demonstrated, with $N = 2,748$ raters and 462 stimuli, that AI beauty-filter modifications shift downstream perceptions of intelligence, trustworthiness, and warmth the classical halo effect, mediated by generative models. The visual channel is not bias-neutral, which is precisely why an image-space probe of SCM is needed in addition to text-space probes.

III. METHODOLOGY

A. Probabilistic Graphical Model

We model each social group g as the source of a latent 2D *stereotype coordinate* $s_g = (W_g, C_g)^\top$ that captures its position in SCM space. Conditional on s_g , a prompt $y_g = \text{“a portrait of a } [g], \text{ high quality”}$ is fed to SD3, producing n images $x_{g,1}, \dots, x_{g,n}$ which are then encoded into high-dimensional embeddings $z_{g,i} \in \mathbb{R}^d$. We make the standard Gaussian assumption per group:

$$z_{g,i} | \mu_g, \Sigma_g \sim \mathcal{N}(\mu_g, \Sigma_g), \quad i = 1, \dots, n. \quad (1)$$

The latent SCM coordinates are recovered from μ_g via projection onto lexicon-defined unit axes $v_W, v_C \in \mathbb{R}^d$:

$$\hat{W}_g = \mu_g^\top v_W, \quad \hat{C}_g = \mu_g^\top v_C. \quad (2)$$

We then relate the model’s coordinates $(\hat{W}_g, \hat{C}_g)$ to the human ground-truth coordinates (W_g^H, C_g^H) through a Bayesian errors-in-variables regression (§III-F), which is the formally Bayesian component of the PGM. The plate diagram is given in Fig. 1.

B. Ground-Truth Data

We use two complementary human-rated sources: (i) the Fiske et al. [1] canonical identity-group set (e.g. *elderly people*, *welfare recipients*, *housewives*), rated on a 5-point Likert scale by ≈ 571 U.S. raters; and (ii) the He et al. [2] occupational SCM ratings of 60 U.S. occupations on a z -scored 7-point scale. The two sources are unified into a single ground-truth file of 66 groups by converting both to a common z -scored (W, C) frame.

C. Synthetic Corpus Generation

We generate 100 images per group using Stable Diffusion 3 Medium (stabilityai/stable-diffusion-3-medium), fixed seeds $\{0, \dots, 99\}$, 30 inference steps, guidance scale 7.0, 1024×1024 resolution. The single template “a portrait of a $[g]$, high quality” was chosen to match prior occupation-bias audits [10] and to keep variation attributable to g . Total corpus: $66 \times 100 = 6,600$ images.




Fig. 1. Plate diagram of the proposed PGM. Each social group g has a latent SCM coordinate $s_g = (W_g, C_g)$ that drives both the human rating s_g^H and the text-to-image generation pipeline (prompt y_g , denoised embedding $z_{g,i}$ from shared SD3 parameters θ_{SD3}). The model coordinate $\hat{s}_g$ is recovered by projecting the per-group embedding mean μ_g onto lexicon-defined warmth and competence axes (Eq. 2). Human and model coordinates are conditionally independent noisy views of s_g , related through the errors-in-variables regression of §III-F.

D. Lexicon-Defined Axes

Following the bipolar-axis construction of Caliskan et al. [9] and Nicolas et al. [4], warmth and competence axes are built by averaging text-encoder embeddings of high-loading positive-pole and negative-pole adjectives and taking the unit-normalized difference:

$$v_W = \frac{\bar{e}^{W+} - \bar{e}^{W-}}{\|\bar{e}^{W+} - \bar{e}^{W-}\|_2}, \quad v_C = \frac{\bar{e}^{C+} - \bar{e}^{C-}}{\|\bar{e}^{C+} - \bar{e}^{C-}\|_2}. \quad (3)$$

Positive-warmth anchors: *friendly, warm, kind, sincere, trustworthy, generous, caring*; negative-warmth: *cold, hostile, unfriendly, unkind, untrustworthy, selfish*. Positive-competence: *intelligent, capable, skilled, competent, efficient, ambitious*; negative-competence: *incompetent, unskilled, incapable, ineffective, lazy*.

E. Dual-Probe Design

The novelty of the methodology relative to text-only SCM probes [6] lies in probing both sides of the diffusion stack:

Text-side probe. The same axis construction (Eq. 3) is applied to each of SD3’s three text encoders: OpenAI CLIP ViT-L/14, OpenCLIP ViT-bigG/14, and Google T5-XXL. Each group name $[g]$ is embedded and projected; correlations with ground truth quantify how much SCM structure each encoder carries.

Image-side probe. Generated images are embedded with (a) CLIP ViT-L/14 (image tower), (b) OpenCLIP bigG (image tower), and (c) DINOv2 ViT-L/14. Axes for (a) and (b) come from the corresponding text tower; for DINOv2, which has no

text tower, we cannot define text-aligned axes and instead use a 4-way logistic-regression quadrant classifier as an axis-free probe. The DINOv2 result is critical for falsifying the CLIP-circularity concern that “CLIP probes find CLIP biases.”

F. Bayesian Errors-in-Variables Regression

Because each $\hat{W}_g$ is itself an estimate from n images, treating it as an exact regressor inflates apparent slope precision. We instead specify

$$\mu_g^{\text{lat}} \sim \mathcal{N}(0, 1) \quad (4)$$

$$\hat{W}_g \mid \mu_g^{\text{lat}} \sim \mathcal{N}(\mu_g^{\text{lat}}, \text{SEM}_g) \quad (5)$$

$$W_g^H \mid \mu_g^{\text{lat}} \sim \mathcal{N}(\alpha + \beta_W \mu_g^{\text{lat}}, \tau), \quad (6)$$

where $\text{SEM}_g = \sqrt{(v_W^\top \Sigma_g v_W)/n}$ is the group-level standard error of the model coordinate. A parallel specification yields β_C . Posteriors are computed with the No-U-Turn Sampler in PyMC (4 chains, 1000 warmup, 1000 draws). The 94% highest density interval (HDI) is reported throughout.

G. Per-Quadrant Kullback–Leibler Divergence

For each SCM quadrant q , let P_q and Q_q be 2D Gaussian fits to the model and human (W, C) coordinates of all groups whose ground-truth quadrant is q . We compute the symmetric KL:

$$\text{KL}_{\text{sym}}(P_q \parallel Q_q) = \frac{1}{2} [\text{KL}(P_q \parallel Q_q) + \text{KL}(Q_q \parallel P_q)]. \quad (7)$$

This pinpoints *which* region of SCM space the model distorts most.

H. Procrustes Alignment and Residuals

To allow for global rotation/reflection/scaling, we compute the optimal Procrustes transform from model to human coordinates and report the residual disparity statistic with a permutation null (1,000 label shuffles). Per-group residuals (Euclidean distance between model and human coordinates after alignment) identify the worst-fit groups for qualitative inspection.

IV. IMPLEMENTATION

A. Pipeline Overview

The full pipeline runs as six modular scripts and consumes ≈ 11 GPU-hours on a single A10G:

- 1) `01_clean_groundtruth.py` merge Fiske and He sources into a single (W, C) ground-truth file, z -score each source independently before union.
- 2) `02_generate_corpus.py` SD3 generation on AWS SageMaker, 100 images/group, deterministic seeds, metadata logging.
- 3) `03_text_side_probe.py` text-axis construction and group-name projections for CLIP-L, bigG, and T5-XXL.
- 4) `04_image_side_probe.py` batch image encoding via CLIP and bigG image towers and via DINOv2; per-group μ_g, Σ_g fits with Ledoit–Wolf shrinkage.

TABLE I
TEXT-ENCODER PROBE (PEARSON r VS. HUMAN RATINGS, $n=66$).

Encoder	r_W	r_C
CLIP ViT-L/14	0.18	0.21
OpenCLIP ViT-bigG/14	0.07	0.49
T5-XXL	0.35	0.51

- 5) 05_pgm_analyses.py Bayesian regression, KL, Procrustes, residuals.
- 6) 06_make_figures.py final figures.

All artifacts are cached to .npz, allowing analyses to be rerun in seconds without re-embedding.

B. Environment and Dependencies

PyTorch 2.3, diffusers 0.30, transformers 4.45, open_clip_torch 2.26, pymc 5.16, scikit-learn 1.5, scipy 1.13. Generation: 1024×1024, 30 steps, guidance 7.0, fp16. Embedding: CLIP-L and bigG via open_clip; DINOv2 via the official Meta checkpoint. All code is released on GitHub (see §IV-D).

C. Reproducibility

Image seeds are fixed at $\{0, \dots, 99\}$ per group; PyMC sampling uses random_seed=42; the full 6,600-image dataset is released on Hugging Face Hub at Shamima/sd3-medium-scm-corpus for verification and reuse.

D. GitHub Repository

Code, figures, lexicons, and a README reproducing every result are available at:

[https://github.com/\[USERNAME\]/scm-probe-sd3](https://github.com/[USERNAME]/scm-probe-sd3)

V. RESULTS AND ANALYSIS

A. The Internalized SCM Map

Figure 2 shows all 66 groups in SD3’s (W, C) space under the OpenCLIP bigG image probe, colored by ground-truth quadrant. A coarse SCM organization is visible: red *contempt* dots cluster in the lower-left, blue *admiration* dots in the upper-right. The *envy* (orange) quadrant, however, collapses inward CEOs, lawyers, and financial advisors are placed near the warmth-axis origin rather than in their ground-truth high- C /low- W position. This visual finding is quantified in §V-E.

B. Text-Encoder Probes

Table I reports Pearson correlations between each SD3 text encoder’s group-name projection and human ratings. T5-XXL recovers both axes ($r_W=0.35$, $r_C=0.51$); OpenCLIP bigG recovers competence but not warmth ($r_C=0.49$, $r_W=0.07$); CLIP-L is weak on both. Because SD3 uses all three encoders in parallel, this asymmetry suggests that warmth signal entering the image is predominantly contributed by T5, while competence is multi-sourced.




Fig. 2. SD3’s internalized SCM map under the OpenCLIP-bigG image probe. Each point is one of the 66 groups; color encodes its *ground-truth* SCM quadrant. Contempt-quadrant groups (red, lower-left) are placed correctly; envy-quadrant groups (orange, ideally lower-right) collapse toward the centre, evidencing systematic compression of high-status outgroups.




Fig. 3. Per-encoder recovery of SCM structure. T5-XXL is the only encoder that carries both warmth and competence; bigG carries competence but not warmth; CLIP-L is weak on both. Because SD3 combines all three at conditioning time, warmth signal entering the image is dominated by T5.

C. Image-Side Probes

Table II reports the analogous correlations on μ_g for the image probes. OpenCLIP-bigG image embeddings recover SCM structure ($r_W=0.48$, $r_C=0.34$); CLIP-L does not. The axis-free DINOv2 4-way quadrant classifier achieves 44% accuracy under 5-fold cross-validation (chance = 25%, binomial $p<0.001$), confirming that SCM-relevant structure exists in the *pixels* themselves, not merely in CLIP’s text-aligned projection of them.

D. Bayesian Errors-in-Variables Regression

The PGM regression (Fig. 5) yields $\beta_W = 0.48$, 94% HDI [0.26, 0.69], $P(\beta_W > 0) \approx 1.0$; and $\beta_C = 0.32$, 94% HDI

TABLE II
IMAGE-SIDE PROBE RESULTS.

Probe	r_W	r_C	4-way acc.
CLIP-L image $\rightarrow$ CLIP-L axes	0.06	0.08	30%
bigG image $\rightarrow$ bigG axes	0.48	0.34	41%
DINOv2 (axis-free)			44%




Fig. 4. Procrustes alignment of SD3 (left) to human SCM (right). Connecting lines emphasize per-group residuals; envy-quadrant groups exhibit the longest displacements.

[0.10, 0.55], $P(\beta_C > 0) \approx 1.0$. Both HDIs exclude zero by a wide margin: SD3’s image-space (W, C) coordinates are positively and significantly aligned with human SCM coordinates after accounting for measurement error in the model coordinates. The Procrustes disparity statistic is 0.81 with permutation $p < 0.001$ (1,000 shuffles), indicating significant joint alignment of the 2D structure even though the r values above are individually modest.

E. Per-Quadrant Distortion

The KL analysis (Table III) localizes the asymmetry hinted at in Fig. 2: *contempt* is reproduced almost faithfully (KL=0.68), while *envy* diverges by a factor of five (KL=3.51). *Paternalism* and *admiration* fall in between. The envy quadrant houses precisely the groups that classical SCM theory characterizes as “competent but cold” high-status outgroups: CEOs, lawyers, financial advisors, politicians. SD3 does not place them where humans do.

F. Worst-Fit Residuals

Figure 6 ranks the 10 groups with the largest post-Procrustes Euclidean error. Eight of the top ten are envy-quadrant occupations (CEO, lawyer, financial advisor, politician, news anchor, salesperson, computer programmer, manager); the remaining two are paternalism-quadrant identity groups whose ground-truth warmth ratings are extreme. The residual structure is not random: SD3 systematically pulls high-status competent-cold occupations toward neutral, consistent with a hypothesized “aspirational portraiture” bias in T2I training data. LinkedIn-style headshots of lawyers and CEOs are uniformly warm-toned and smiling.




Fig. 5. Posterior densities for the warmth and competence regression slopes β_W and β_C from the errors-in-variables PGM. Vertical dashed line at $\beta = 0$; shaded 94% HDIs both exclude zero.

TABLE III
PER-QUADRANT SYMMETRIC KL DIVERGENCE (LOWER = BETTER FIT).

Quadrant	KL _{sym}
Contempt (cold, incompetent)	0.68
Paternalism (warm, incompetent)	1.66
Admiration (warm, competent)	1.91
Envy (cold, competent)	3.51

G. Summary of Findings

SD3 internalizes the SCM topology to a statistically robust but quantitatively partial degree. The warmth dimension weakly carried by CLIP and bigG text encoders — emerges in image space, plausibly contributed by T5 conditioning. The asymmetry between contempt and envy quadrants is the central empirical contribution: *frequency* audits miss it because high-status occupations are not numerically under-represented; they are geometrically misplaced.

VI. USEFULNESS

The methodology and findings are useful in three settings.

A. For AI Auditing

Standard T2I bias audits report *marginal* statistics: gender ratio per occupation, racial composition per prompt. These statistics do not detect the envy-compression effect we report: a CEO image set can be demographically balanced while still failing to convey the “competent-cold” visual signature that humans associate with high-status outgroups. Our SCM-aligned probe offers a complementary, *geometric* audit lens that surfaces this class of failure. Auditors can run the probe (six scripts, one GPU-day) on any new T2I model and obtain a quadrant-level distortion report.

B. For Probabilistic Modeling

The errors-in-variables PGM is, to our knowledge, the first explicit Bayesian treatment of T2I representational alignment with a human-rated psychometric construct. The framework generalizes to other 2D social-cognition constructs (e.g. agency–communion, ideology–engagement maps) and to other generative models. The PGM framing also clarifies what existing CLIP-style probes *assume* namely, that the encoder is a noise-free observation channel and corrects that assumption.




Fig. 6. Top-10 worst-fit groups after Procrustes alignment. Eight of ten belong to the envy quadrant, reinforcing the quadrant-level KL finding.

C. For Social Psychology

The portrait corpus is itself a research artifact: 6,600 SCM-annotated images of 66 groups, generated with a known production T2I model. Social-cognition researchers can use the corpus directly as stimuli in human IAT or feature-rating experiments, contributing data on *generated* stimulus perception that complements the longstanding photographic SCM literature [3].

VII. LIMITATIONS AND FUTURE WORK

A. Limitations

a) Single-model scope: All results concern SD3 Medium. We do not yet know whether the envy-quadrant compression generalizes to SDXL, SD3.5, FLUX, or commercial systems such as DALL-E 3 or Imagen.

b) Lexicon dependence: Warmth and competence axes depend on the chosen anchor adjectives (Eq. 3). We used a ≈ 7 -word lexicon per pole drawn from Nicolas et al. [4]; a sensitivity analysis across two alternative lexicons (omitted for space) confirms the β estimates fall within the original 94% HDI, but the issue cannot be fully removed in a lexicon-defined probe.

c) Anglocentric ground truth: Both Fiske [1] and He [2] ratings are U.S. samples. SCM has been replicated cross-culturally [5], but global SCM mappings would require rebuilding the ground-truth file.

d) No human image-ratings: We compare SD3 image projections to humans rating *group labels*, not to humans rating the *generated portraits* themselves. A human study rating ~ 200 portraits on warmth and competence would let us distinguish “SD3 fails to encode the stereotype” from “SD3 encodes it but CLIP/DINOv2 cannot read it.”

e) Prompt sensitivity: We use a single template. Auditing under templates that condition on context (“a CEO at work,”

“a CEO at home”) would localize where in semantic context the envy distortion is strongest.

B. Future Work

a) Causal mechanism: The text-encoder ablation suggests warmth enters via T5 and competence via bigG. Causal interventions swapping or perturbing the T5 conditioning signal during denoising would test this directly.

b) Cross-model replication: Running the same probe across SD3.5, SDXL, FLUX, and a closed-source API model would establish whether envy-compression is a property of one model or of the diffusion paradigm at large.

c) Mitigation: Once the mechanism is known, prompt-level or LoRA-level interventions could be designed to restore envy-quadrant geometry without inducing demographic side-effects. This is the constructive flip side of the audit and the natural next step.

d) Beyond SCM: The same dual-probe + Bayesian-EIV framework applies to any 2D psychometric construct grounded in a validated lexicon, including agency–experience [17] and political ideology maps. SCM is one instance of a broader class of structural-bias probes that the methodology enables.

ACKNOWLEDGMENTS

We thank the instructors of the Probabilistic Graphical Models course for guidance, and the maintainers of the Fiske and He datasets for releasing their ratings publicly.

REFERENCES

- [1] S. T. Fiske, A. J. C. Cuddy, P. Glick, and J. Xu, “A model of (often mixed) stereotype content: competence and warmth respectively follow from perceived status and competition,” *J. Pers. Soc. Psychol.*, vol. 82, no. 6, pp. 878–902, 2002.
- [2] J. C. He, S. K. Kang, K. Tse, and S. X. Toh, “Stereotypes at work: Occupational stereotypes predict race and gender segregation in the workforce,” *J. Vocational Behavior*, vol. 115, 103318, 2019.
- [3] R. Imhoff and J. Woelki, “Warmth and competence in your face! Visual encoding of stereotype content,” *Frontiers in Psychology*, vol. 4, p. 386, 2014.
- [4] G. Nicolas, X. Bai, and S. T. Fiske, “A spontaneous stereotype content model: taxonomy, properties, and prediction,” *Eur. J. Soc. Psychol.*, vol. 51, no. 1, pp. 178–196, 2021.
- [5] A. J. C. Cuddy et al., “Stereotype content model across cultures: Towards universal similarities and some differences,” *British J. Soc. Psychol.*, vol. 48, pp. 1–33, 2009.
- [6] K. C. Fraser, S. Kiritchenko, and I. Nejadgholi, “How does stereotype content differ across data sources?” in *Proc. *SEM*, 2024, pp. 18–34.
- [7] G. Nicolas and A. Caliskan, “Construct differentiability in the stereotype content model,” University of Pennsylvania, Working paper, 2024.
- [8] “An exploratory semantic analysis of age-related stereotypes in OpenAI’s GPT-4o model,” *The Gerontologist*, 2025.
- [9] A. Caliskan, J. J. Bryson, and A. Narayanan, “Semantics derived automatically from language corpora contain human-like biases,” *Science*, vol. 356, no. 6334, pp. 183–186, 2017.
- [10] F. Bianchi et al., “Easily accessible text-to-image generation amplifies demographic stereotypes at large scale,” in *Proc. FAccT*, 2023.
- [11] S. Ghosh and A. Caliskan, “Person, human, neither: the dehumanization potential of automated image tagging,” in *Proc. AIES*, 2024.
- [12] V. Hofmann, P. R. Kalluri, D. Jurafsky, and S. King, “AI generates covertly racist decisions about people based on their dialect,” *Nature*, vol. 633, no. 8028, pp. 147–154, 2024.
- [13] J. Ho, A. Jain, and P. Abbeel, “Denoising diffusion probabilistic models,” in *Proc. NeurIPS*, 2020.
- [14] R. Rombach, A. Blattmann, D. Lorenz, P. Esser, and B. Ommer, “High-resolution image synthesis with latent diffusion models,” in *Proc. CVPR*, 2022.

- [15] G. Gala, C. de Campos, A. Vergari, and E. Quaeghebeur, "Scaling continuous latent variable models as probabilistic integral circuits," in *Proc. NeurIPS*, 2024.
- [16] A. Gulati *et al.*, "What is beautiful is still good: the attractiveness halo effect in the era of beauty filters," *R. Soc. Open Sci.*, vol. 11, 240882, 2024.
- [17] H. M. Gray, K. Gray, and D. M. Wegner, "Dimensions of mind perception," *Science*, vol. 315, no. 5812, p. 619, 2007.